

????

Matthew Paletta

# Docker ???

Matthew Paletta

# Docker + gRPC ??


Matthew Paletta

# Docker + gRPC + Redis ?

Matthew Paletta

# **Docker + gRPC + Redis + DevOps**

Matthew Paletta



**“Everything is figureoutable”**

# Disclaimer

Buzzwords - feel free to ask me!



# Ground Zero

# Ground Zero

- This is our starting spot, based on what we've seen in the labs
-

# Ground Zero


- One process

# Ground Zero

- One process
- One machine


# Ground Zero

- One process
- One machine
- Multi-threaded?




- Here's our world view. We only think about our code

# Goal



- We'd like to put it on 2 machines!
- Maybe this is a web server for our website, and we want to handle more requests. We want to handle more requests, or have it fault-tolerant, so our website doesn't go down.


# Goal



- And once we do 2, we want 3



# Goal




## Scaling! (distributing)

- And when we do that, we want to use other peoples machines (cloud)
- And we want to scale and distribute to increase the amount of data we can process


# Containers

- What is this container thing -

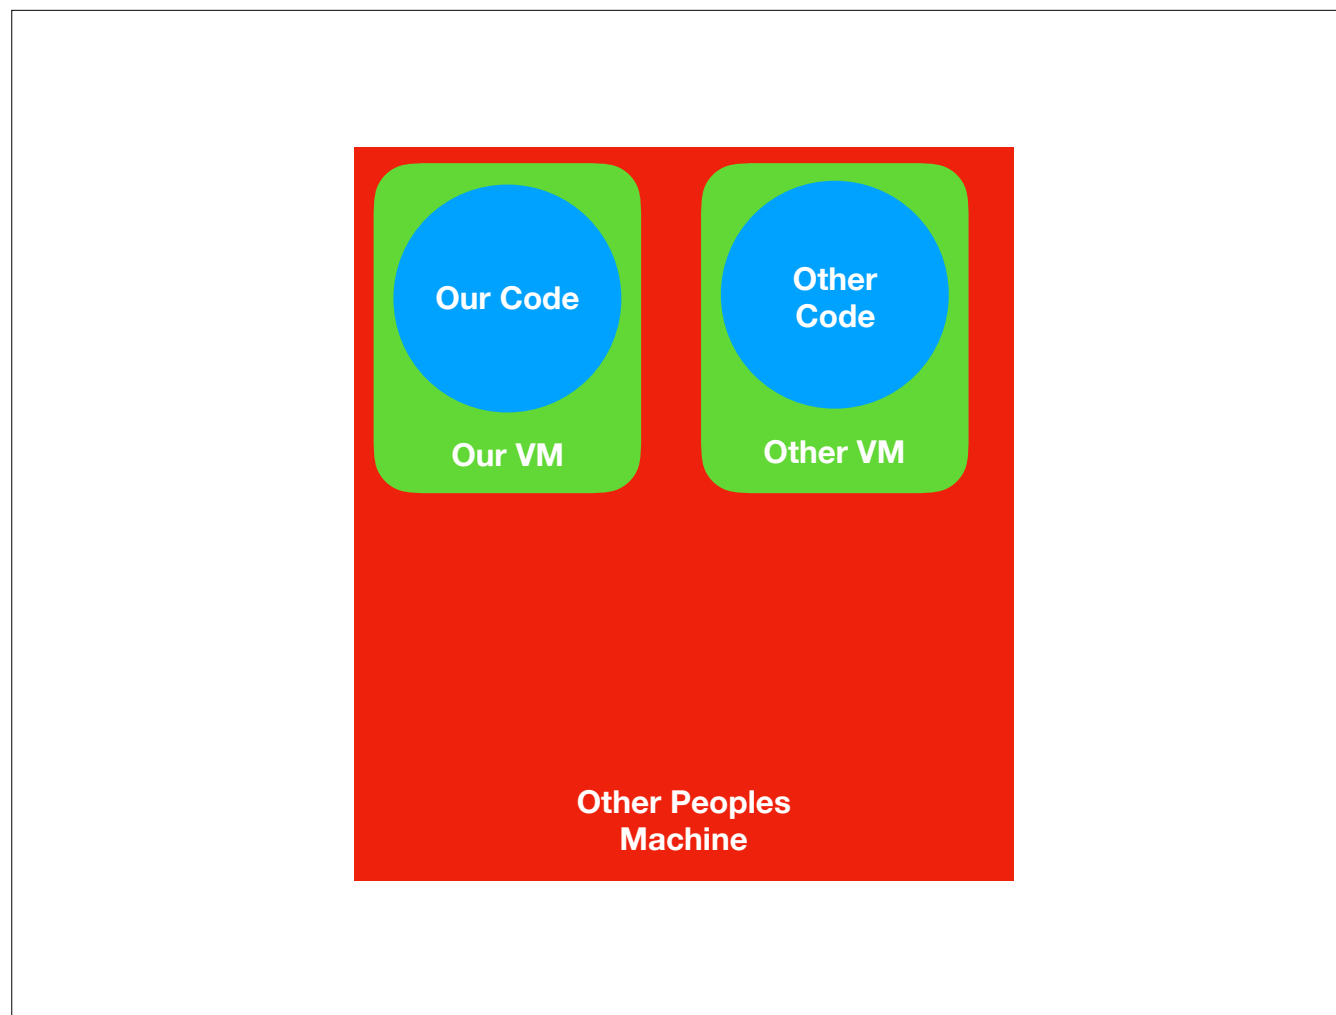


- Lets start again with our code running on our machine





Let's say we share that machine with other people  
This is what it actually looks like



- Probably there's other VMs running on the machine too.

# Virtual Machines vs. Containers

- We could look at this differently

**Our Machine**



Infrastructure

**Our Machine**

Host Operating System

Infrastructure


**Our Machine**

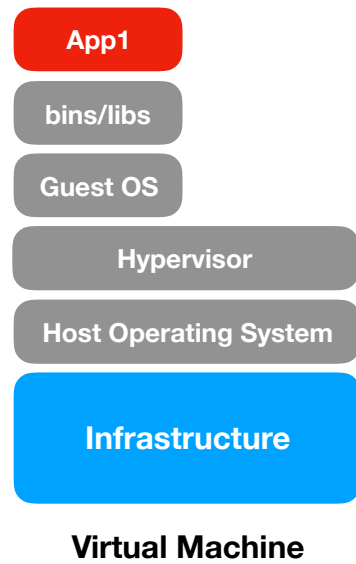
bins/libs

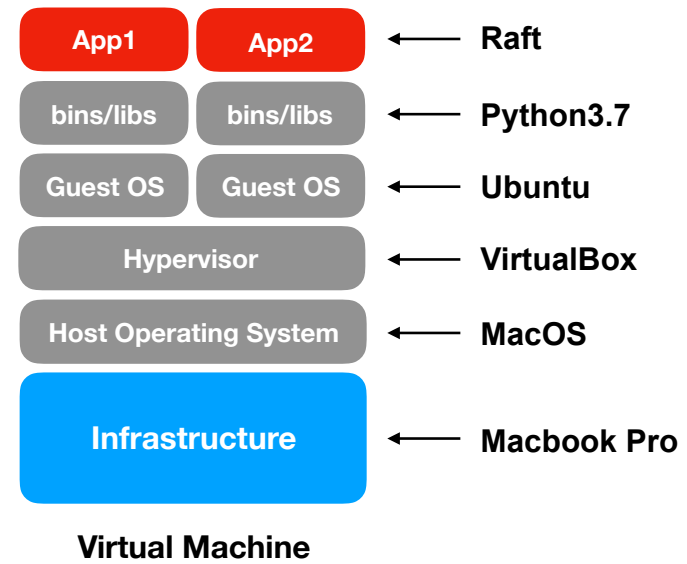
Host Operating System

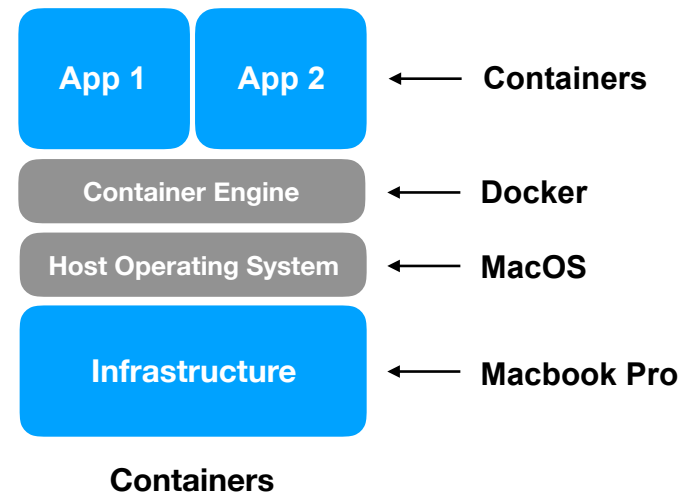
Infrastructure

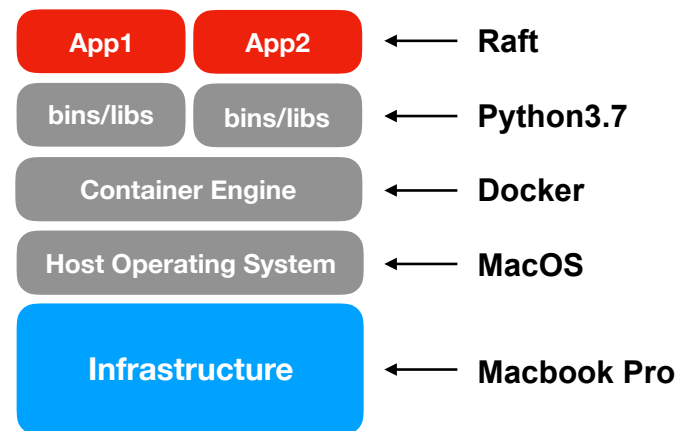
**Our Machine**





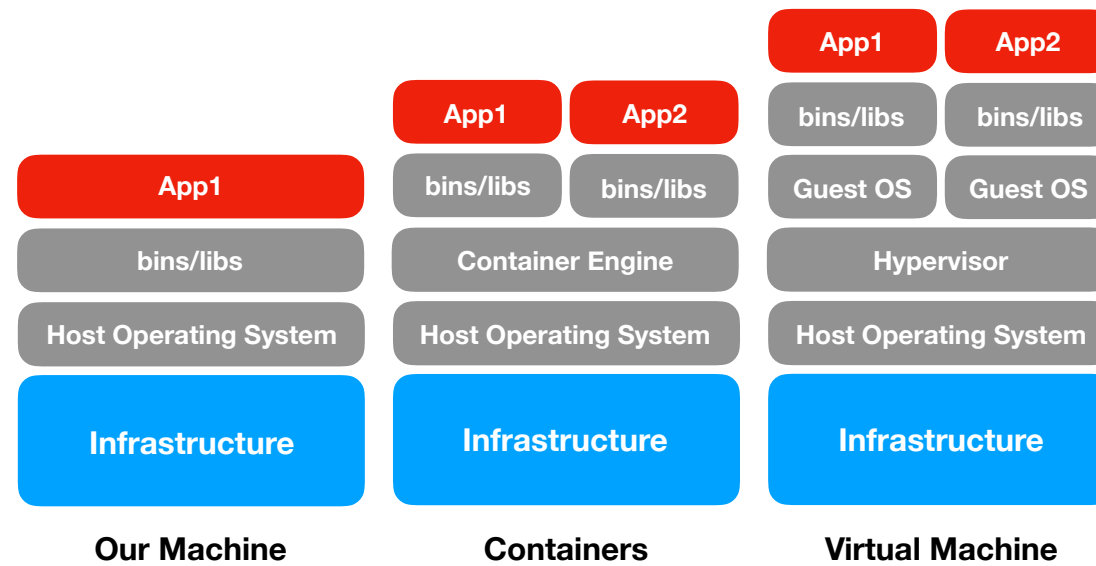






**Containers**





- Containers give us better performance - close to bare-metal, but each app can have each bins/libs
- Containers are distributable environment
- Ask why containers

# Containers

- Better performance compared to VMs
- Only consume the resources they are actually using (disk & memory)
- Maintains per-app isolation & libraries
- Easy to distribute and scale with other developers or cloud deployments
- Faster startup time - don't have to wait for the entire OS to boot!
- Infrastructure as code - (devops topic)

# Docker

# Dockerfiles

# Our Example Project

```
project_root/  
|  our_example_package/  
|  —> main.py  
|  —> resources/  
|      | —> .....  
| Dockerfile  
| README.md  
| requirements.txt  
| .....
```

- This is our example project structure
- Nothing too special going on here

# Dockerfile

- Here's the Dockerfile (stepped through)
- Introduce a 'layer'
- There are other 'FROM' images - we'll come back
- Ask if there's any questions about any of the lines

# Dockerfile

```
FROM python:3.7
```

# Dockerfile

```
FROM python:3.7
```

```
ADD requirements.txt /requirements.txt
```



# Dockerfile

```
FROM python:3.7
```

```
ADD requirements.txt /requirements.txt
```

```
RUN pip3 install -r requirements.txt
```

# Dockerfile

FROM python:3.7

ADD requirements.txt /requirements.txt

RUN pip3 install -r requirements.txt

WORKDIR basic\_example\_project

# Dockerfile

FROM python:3.7

ADD requirements.txt /requirements.txt

RUN pip3 install -r requirements.txt

WORKDIR basic\_example\_project

ADD our\_example\_package .

# Dockerfile

FROM python:3.7

ADD requirements.txt /requirements.txt

RUN pip3 install -r requirements.txt

WORKDIR basic\_example\_project

ADD our\_example\_package .

ENTRYPOINT ["python3", "main.py"]

# Dockerfile

```
FROM python:3.7
```

```
ADD requirements.txt /requirements.txt
```

```
RUN pip3 install -r requirements.txt
```

```
WORKDIR basic_example_project
```

```
ADD our_example_package .
```


```
ENTRYPOINT ["python3", "main.py"]
```

WORKDIR - the directory in the `container!` we are running subsequent commands (until further notice)

ENTRYPOINT - command to run when starting the container


CMD - appended to entrypoint, provides arguments

# ENTRYPOINT vs. CMD




our\_image  
(latest)

# ENTRYPOINT vs. CMD



**WORKDIR /**  
**ENTRYPOINT: ["ls"]**

# ENTRYPOINT vs. CMD




**WORKDIR /**  
**ENTRYPOINT: ["ls"]**

`docker run our_image:latest`

**calls 'ls /'**



# ENTRYPOINT vs. CMD



**WORKDIR /**  
**ENTRYPOINT: ["ls"]**

<code>docker run our_image:latest</code>	<b>calls 'ls /'</b>
<code>docker run our_image:latest example_dir</code>	<b>calls 'ls /example_dir'</b>

# ENTRYPOINT vs. CMD



**WORKDIR /**  
**ENTRYPOINT: ["ls"]**

[illegible]

**Your Turn!**

# Docker Hub

- Allows sharing of docker images
- Images such as `python:3.7` are hosted on Docker Hub
- <https://hub.docker.com>
- Signup for a free account!
- Use `docker login` on your machine

# Building Images

- `docker image build -t <image_name>:<image_tag> <working_dir>`
- e.g. *docker image build -t run\_docker:latest .*
- Can manually specify *Dockerfile* with *-f*
- All files you copy into your docker image must be in, or subdirectories of *working\_dir*

# Docker Labs

- <https://labs.play-with-docker.com>
- Free learning resources in a browser
- Can use one or more machines to experiment with Docker, Docker-Compose & Docker Swarm

# Basic Task

- run an instance of the `redis:latest` image on your machine
- build and run a image/container for the code in `exercises/docker/run_docker/` on Github

# Adv. Task

- search and run an image other than Python or Redis on Hub
- try running it with `-d` - what does that do?
- Look at your current containers with `docker ps`
  - What happens if you run it `--rm`
- Look at your current images with `docker images`
  - What's the difference between `redis:latest` and `redis:alpine`

## Helpful commands:

`docker run <image>`

`docker image build -t <img>:<tag> <work_dir>`


Solutions on Github

# RPC

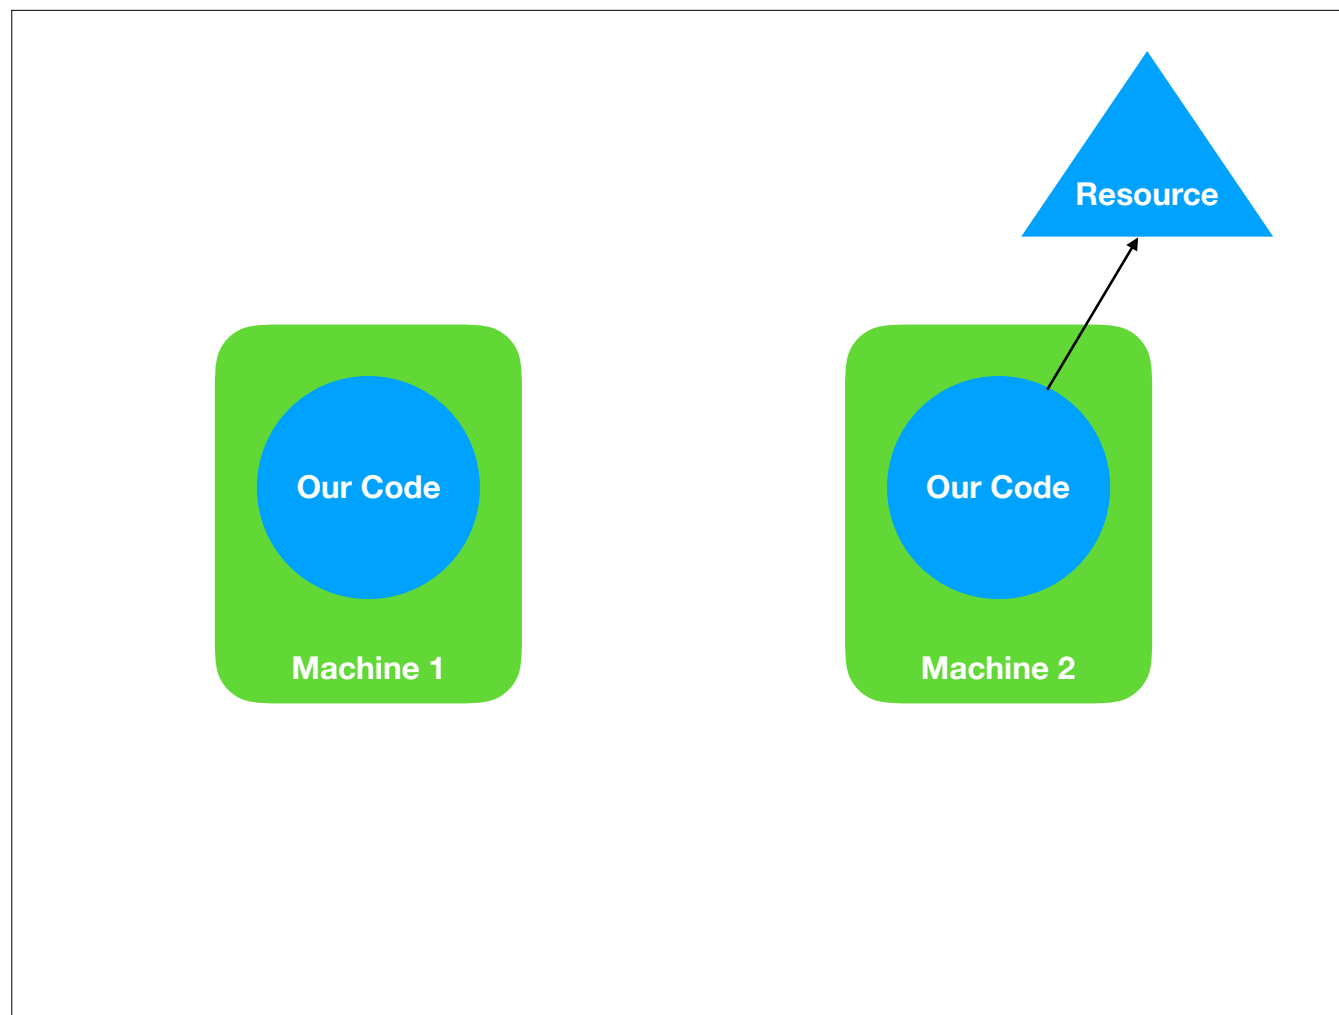
- I want to show you one more Docker thing that will let you run multiple containers and connect them, but first, RPC!

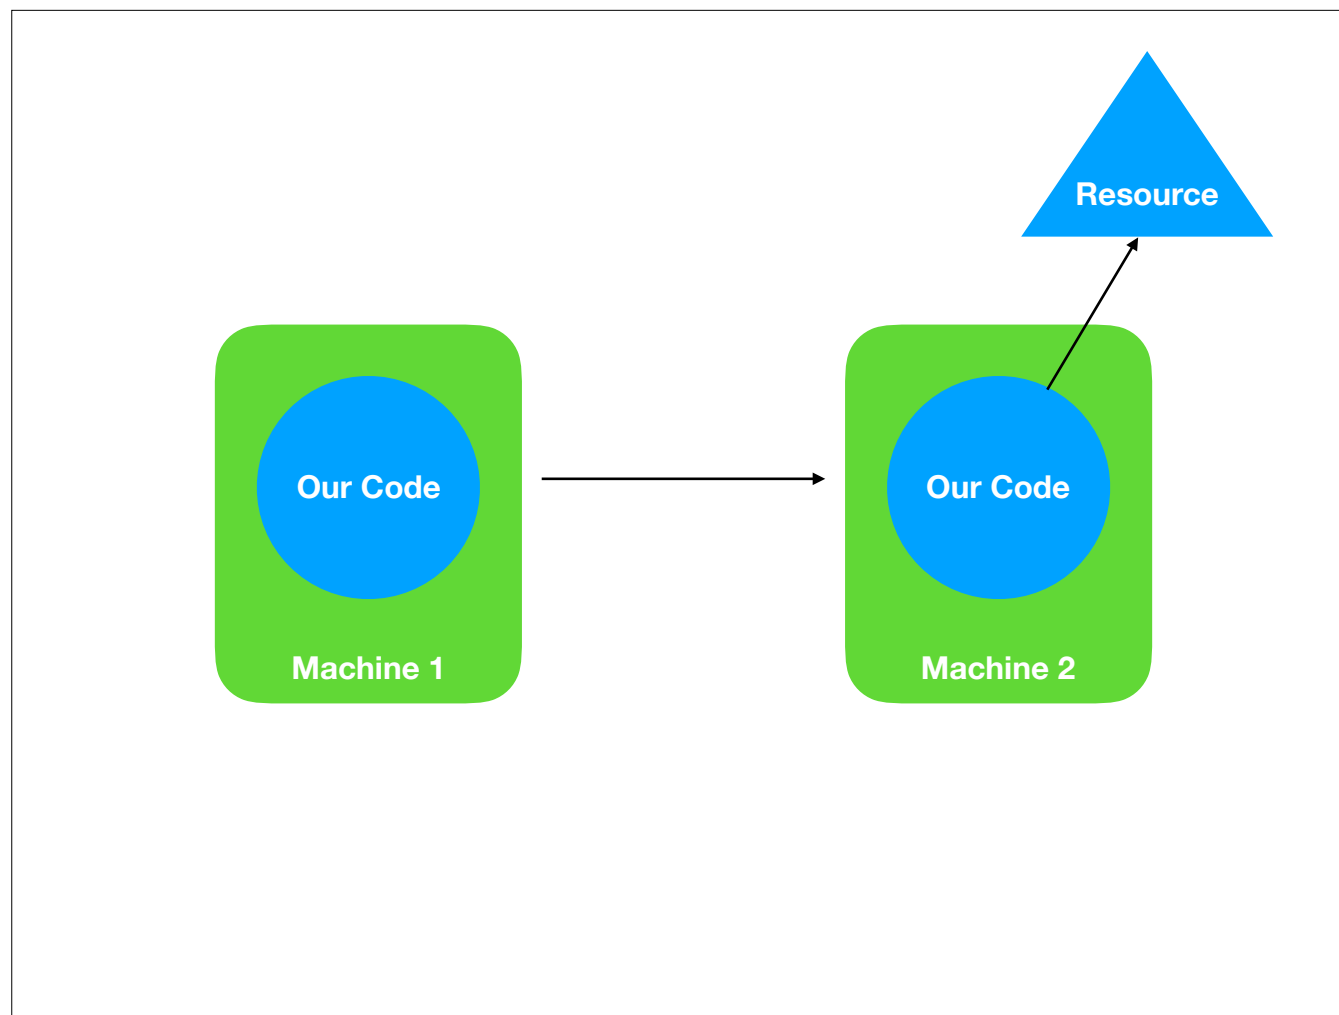


# Remote Procedure Call



Goal - we have 2 systems on the network, we want to run a function from one machine to another  
Maybe the second machine has access to certain data or resources we want





# j(son)RPC

- Consider this thought experiment
- Lets implement a jRPC Protocol

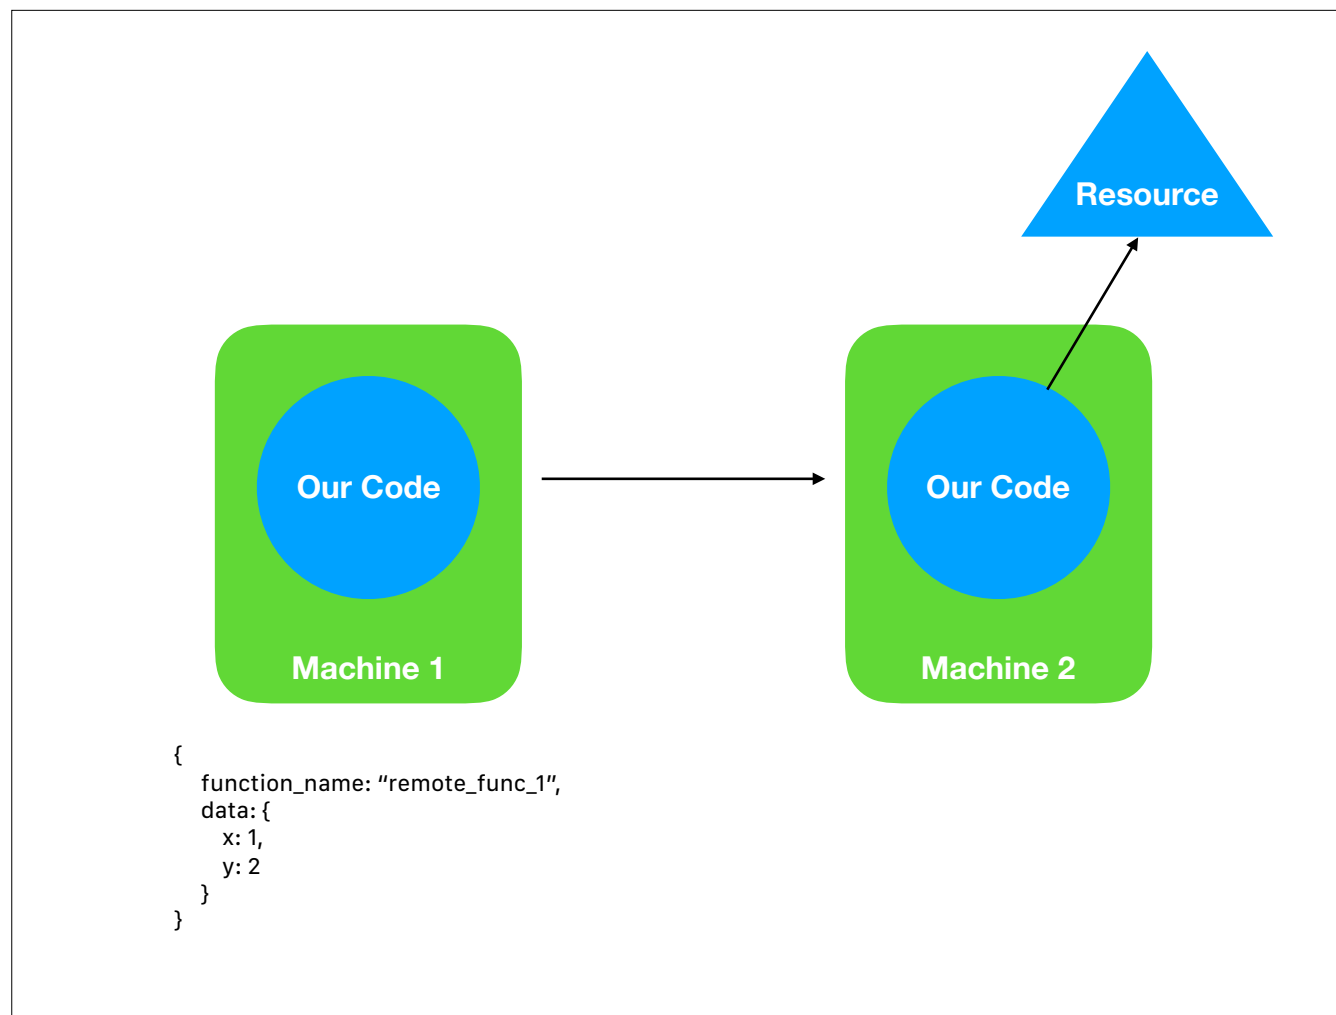
# j(son)RPC

```
// Simple jRPC Schema
{
  function_name: "...",
  data: {
    ...
    ...
  }
}
```

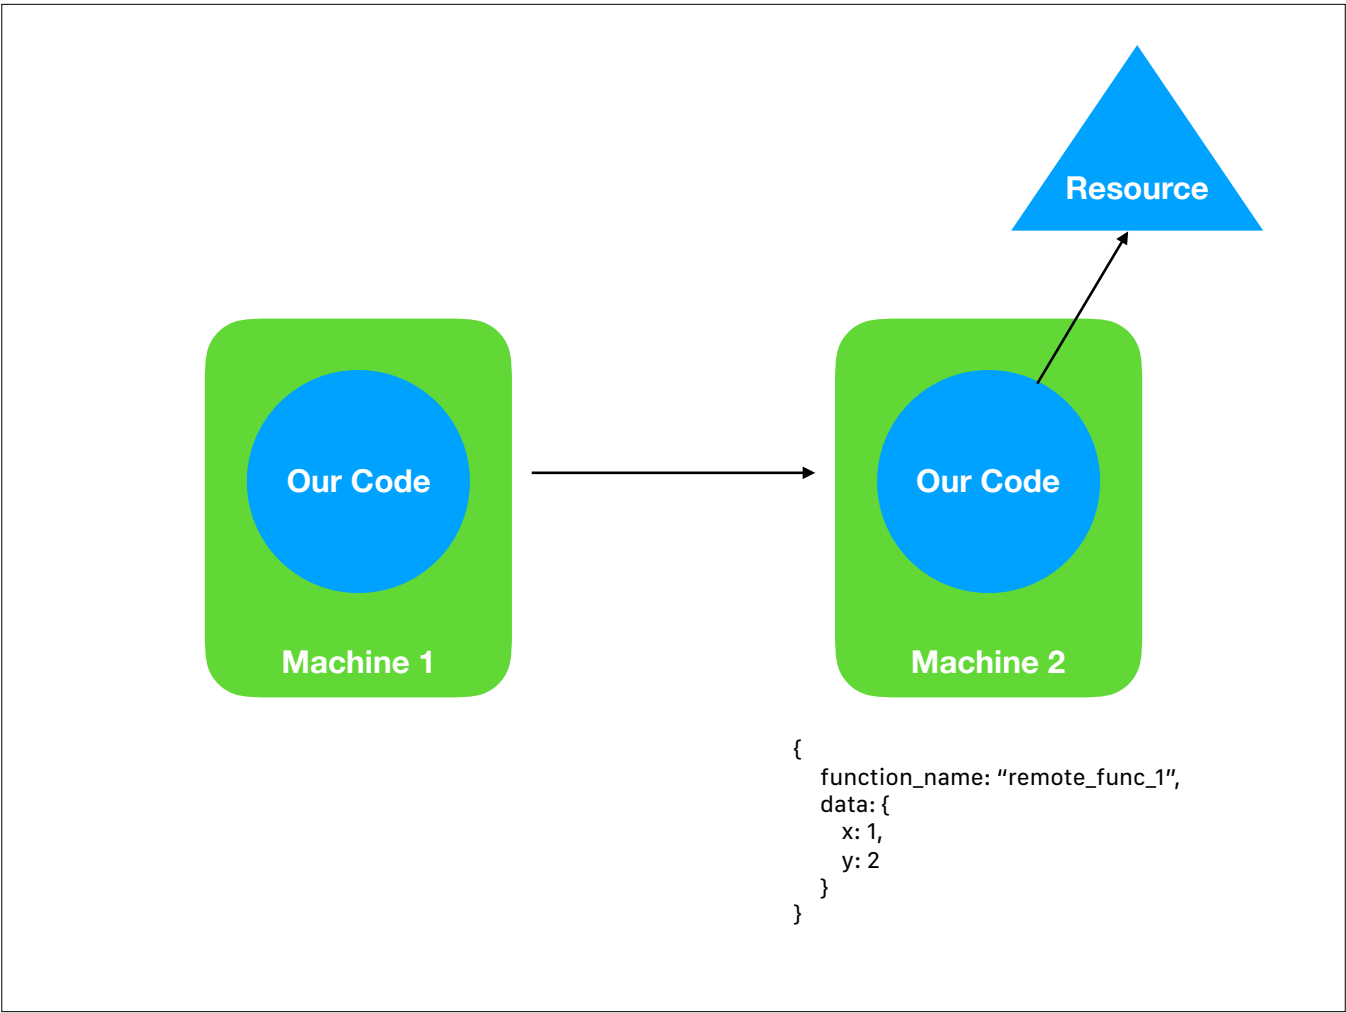
# j(son)RPC

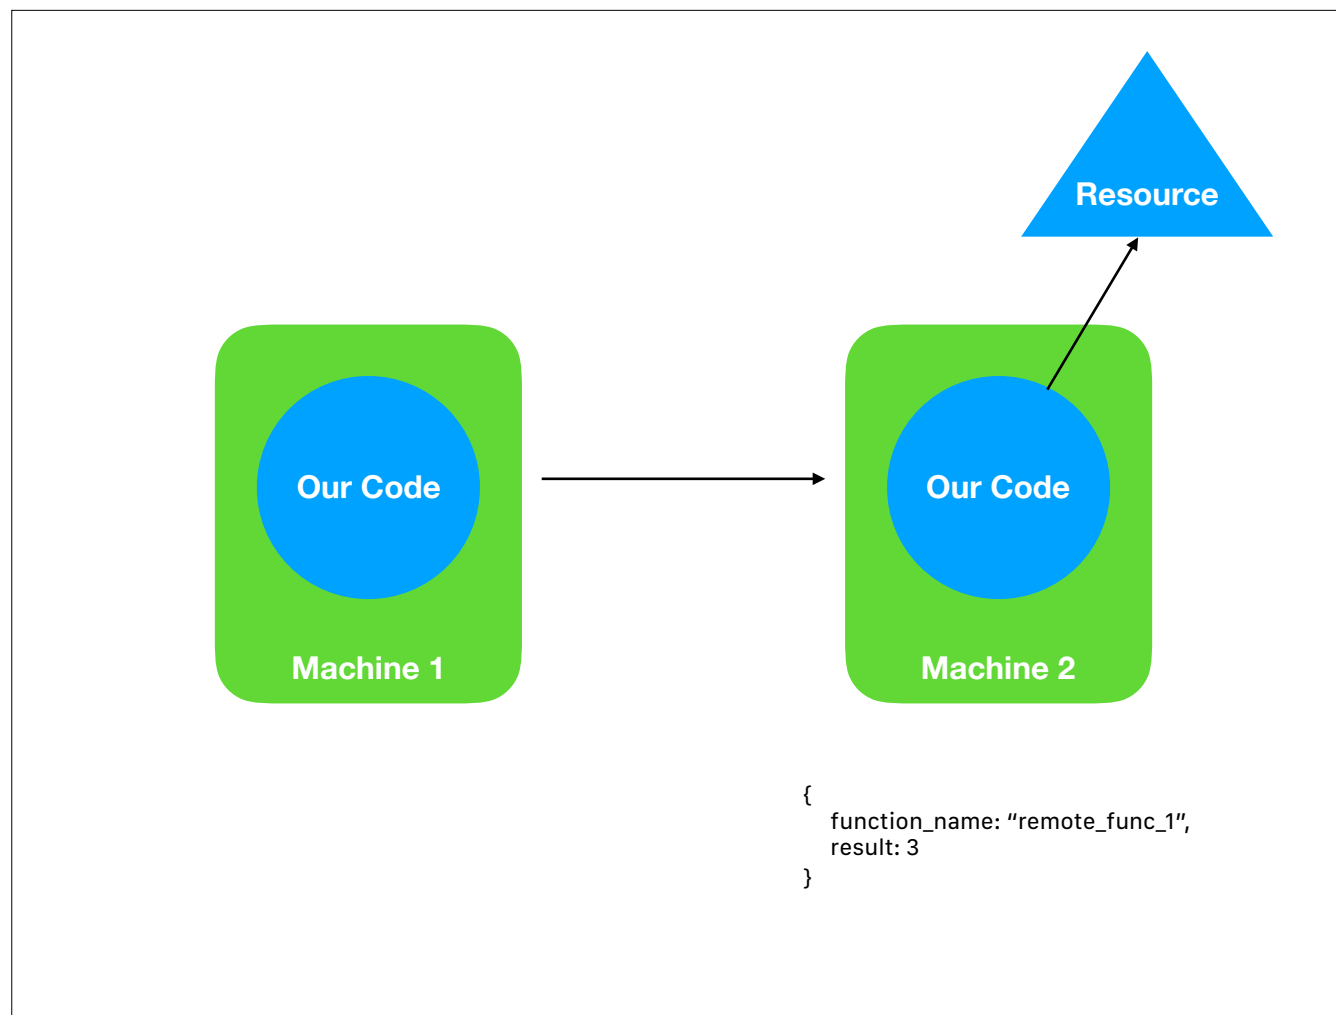
```
// Simple jRPC Schema
{
  function_name: "remote_func_1",
  data: {
    x: 1,
    y: 2
  }
}
```

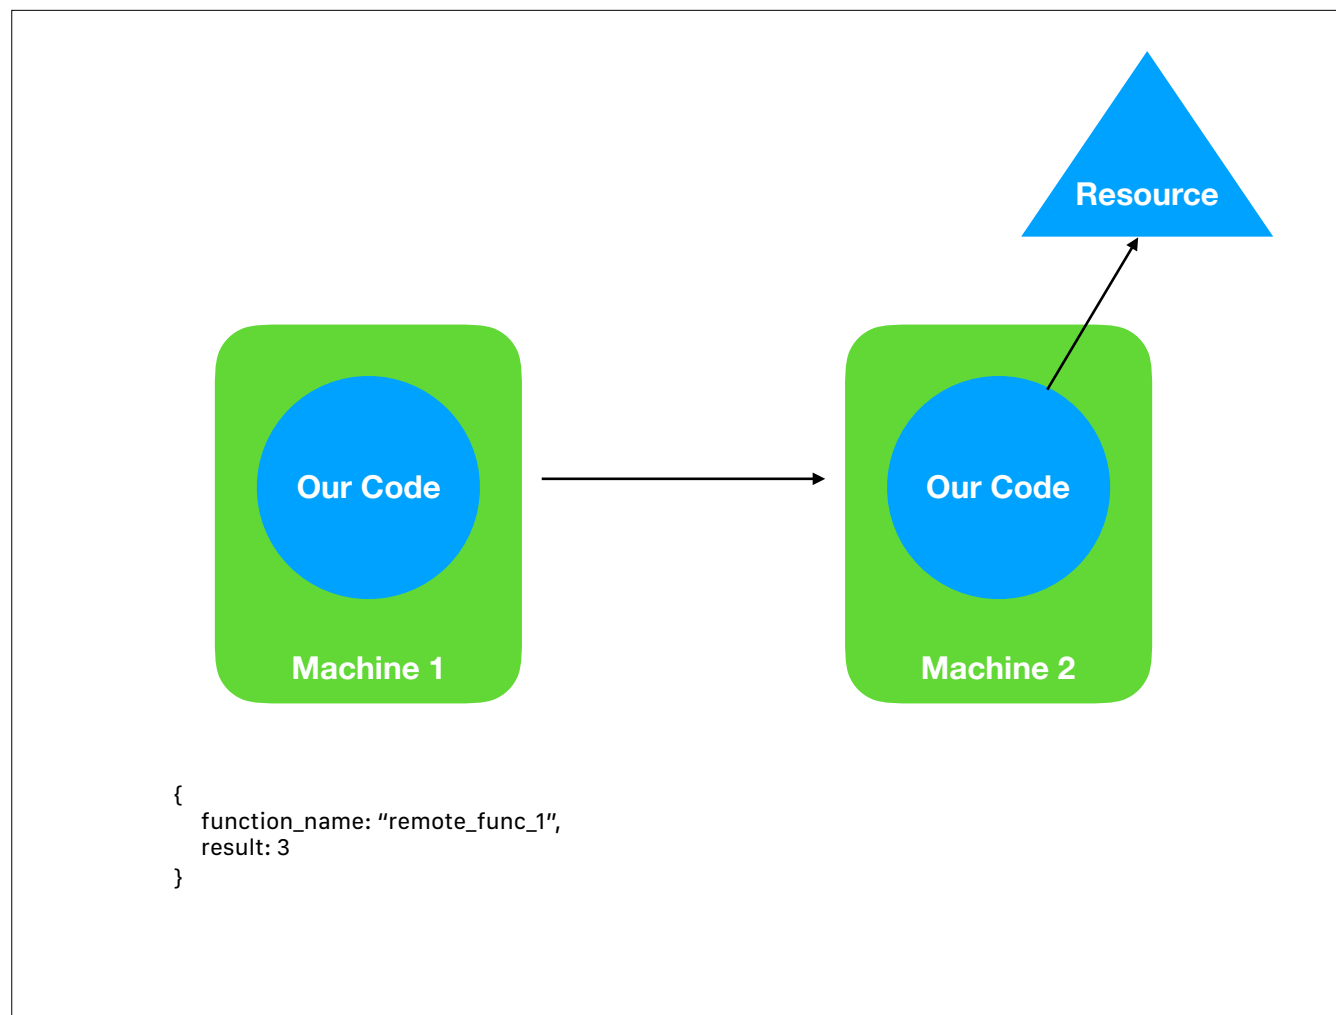
- we could fill in our code with values such as the following











# jRPC Implementation

```
func remote_func_1(int x, int y) { ... }

func getNextMessage(chan) {
    msg = chan.getNext()
    if msg.func_name == "remote_func_1" {
        result = self.remote_func_1(msg.x, msg.y)
        return result.toJSON()
    } else if (...) {
        ...
    } else {
        return FunctionNotFoundError(msg).toJSON()
    }
}
```

Here's how we might implement it

**gRPC**

# gRPC

- Does not send schema (JSON), only binary data (Protobuf)
- Static typing & compiled
- Supports multiple programming languages

basic

advanced

- Synchronous & Asynchronous communication
- Only RPC to support streaming

<https://www.grpc.io/blog/principles/>

- Don't worry if you don't understand the last 2!

[github.com/mattpaletta/distributed-docker-example](https://github.com/mattpaletta/distributed-docker-example)

# Hello gRPC Example

```
hello_grpc/  
| helloworld.proto  
| main_client.py  
| main_server.py  
| requirements.txt  
| Makefile (for your convenience)
```

- Here's our folder structure so you can follow along



# requirements.txt

```
grpcio==1.21.1  
grpcio-tools==1.21.1 #(optional)
```

- Here's our folder structure so you can follow along

# helloworld.proto

```
syntax = "proto3";  
package hellogrpc.helloworld;  
  
message HelloData {  
    int64 x = 1;  
    int64 y = 2;  
    string user = 3;  
}  
  
message GoodbyeData {  
    int64 the_thing = 1;  
}  
  
service HellogRPC {  
    rpc do_the_thing(HelloData) returns (GoodbyeData) {}  
}
```

- I called them different things just so you can follow along. They could be the same types! - just can't be empty

## <run>

```
python3 -m grpc_tools.protoc -I./ \
  --python_out=./ \
  --grpc_python_out=./ \
  helloworld.proto
```

Generates:

```
helloworld_pb2.py
helloworld_pb2_grpc.py
```

`make protos`

\* if you didn't install `grpcio-tools`, call `protoc` directly

# main\_server.py

- The red things are determined by your proto file and will change per implementation
- The orange class is the class we put our implementation in
- I've added type annotations so its easier to follow

# main\_server.py

```
class HellogRPCServiceImpl(HellogRPCServicer):
    def do_the_thing(self, request: HelloData, context: grpc.RpcContext = None):
        x = request.x
        y = request.y
        user = request.user
        print("Got request from: ", user)
        return GoodbyeData(the_thing = x + y)
```

# main\_server.py

```
class HellogRPCServiceImpl(HellogRPCServicer):
    def do_the_thing(self, request: HelloData, context: grpc.RpcContext = None):
        x = request.x
        y = request.y
        user = request.user
        print("Got request from: ", user)
        return GoodbyeData(the_thing = x + y)

if __name__ == "__main__":
    hello_port = 4620
    server = grpc.server(futures.ThreadPoolExecutor(max_workers = 4))
    add_HellogRPCServicer_to_server(servicer = HellogRPCServiceImpl(), server = server)
    server.add_insecure_port('localhost:{0}'.format(hello_port))
    server.start()
    while True:
        sleep(ONE_DAY_IN_SECONDS)
```

# main\_client.py

```
if __name__ == "__main__":  
    channel = grpc.insecure_channel('localhost:4620')  
    hello_stub = HelloRPCStub(channel)  
  
    hello_data = HelloData(x = 1, y = 2, user = "student")  
    response: GoodbyeData = hello_stub.do_the_thing(hello_data)  
    print("Received: " + response.the_thing)
```

- The red things are determined by your proto file and will change per implementation
- The orange class is the class we put our implementation in
- I've added type annotations so its easier to follow

# One more thing

- The red things are determined by your proto file and will change per implementation
- The orange class is the class we put our implementation in
- I've added type annotations so its easier to follow



# One more thing

Python 3.6+ has type annotations! (as you saw)

# One more thing

Python 3.6+ has type annotations! (as you saw)

Package: [mypy-protobuf](#)

# One more thing

Python 3.6+ has type annotations! (as you saw)

Package: **mypy-protobuf**

```
python3 -m grpc_tools.protoc -I./\
  --python_out=./\
  --grpc_python_out=./\
  --mypy_out=./\
  helloworld.proto
```

Generates type-annotations for your IDE!

(example in `hello_grpc_typed`)

**Your turn!**

# Basic Task

- Given a *.proto*, and a server, write the client (my\_client)
- Given a *.proto*, write the server and the client (my\_server)
  - You don't have to lookup the weather! (but you can)
  - Fake the data!

# Adv. Task

- Collaborate with someone, one of you write the *.proto* and the client, the other person write the server!
- Try and explore other data types in Protobufs - repeated & enums

## Helpful commands:

```
python3 -m grpc_tools.protoc -I./ \
--python_out=./ \
--grpc_python_out=./ \
helloworld.proto
```

**docker compose -> swarm**


- We'll start with docker compose, and just introduce swarm

# YAML

- YAML Ain't Markup Language
- Think JSON
  - With comments
  - Uses indentation instead of brackets

[https://docs.ansible.com/ansible/latest/reference\\_appendices/YAMLSyntax.html](https://docs.ansible.com/ansible/latest/reference_appendices/YAMLSyntax.html)

# Hello Compose



**Think of Redis as a Database for now!  
We'll come back to Redis in a few slides**

- we could also use MySQL or Postgres as examples, just Redis is one of the simplest



# Hello Compose Example

```
hello_compose/  
| docker-compose.yml  
| foo/  
|   __main__.py  
|   requirements.txt  
|   Dockerfile
```

- Here's our folder structure so you can follow along

# foo/requirements.txt

redis==3.2.1

# foo/\_\_main\_\_.py

```
import redis
```

```
if __name__ == "__main__":
```

```
    redis_conn = redis.StrictRedis(host = "bar", port = 6379)
```

```
        redis_conn.set(name = "cat", value = 1)
```

```
        redis_conn.set(name = "dog", value = 2)
```

```
        redis_conn.set(name = "mouse", value = 3)
```

```
    result = str(redis_conn.get(name = "cat"))
```

```
    print(result)
```

# docker-compose.yml

```
version: "3"
```

```
services:
```

```
  foo:
```

```
    build: foo
```

```
    links:
```

```
      - bar
```

```
  bar:
```

```
    image: redis:alpine
```

# <run>

*docker-compose up*

- Most basic command
- Runs everything
- Builds and pulls images - as needed

*docker-compose up --build*

- Explicitly rebuilds images - with a cache

*docker-compose up foo*

- Run specific service(s)
- Automatically runs all dependencies

*docker-compose up -d bar*

- Runs service 'bar' in detached mode - advanced

# <run> (advanced list)

```
docker-compose up -d  
docker-compose down  
docker-compose start  
docker-compose stop  
docker-compose build  
docker-compose logs -f db  
docker-compose scale db=4  
docker-compose events  
docker-compose exec db bash
```

# Docker Swarm

- Compose used for testing multi-container on your machine
- Swarm used to deploy to a group of machines
  - Uses docker-compose.yml as configuration
- Alternatives include Kubernetes, and Cloud solutions
- Save for later - feel free to ask me :)

**Your Turn!**



# Basic Task

- Take the server and client you wrote previously, create a `docker-compose.yml` file and run it!

# Adv. Task

- Try adding a Redis cache as part of your compose. (next section of slides)
- Try adding a different database - Postgres, MySQL, etc.

## Helpful commands:

`docker-compose up --build`

[Solutions on Github](#)


# Redis

key/value store, distributed queue, etc.


# Redis

- In-Memory data structure store
- Commonly used as a cache or message queue
- Supports storing:
  - Strings
  - Hashes
  - Lists
  - Sets (sorted & unsorted)
  - Streams

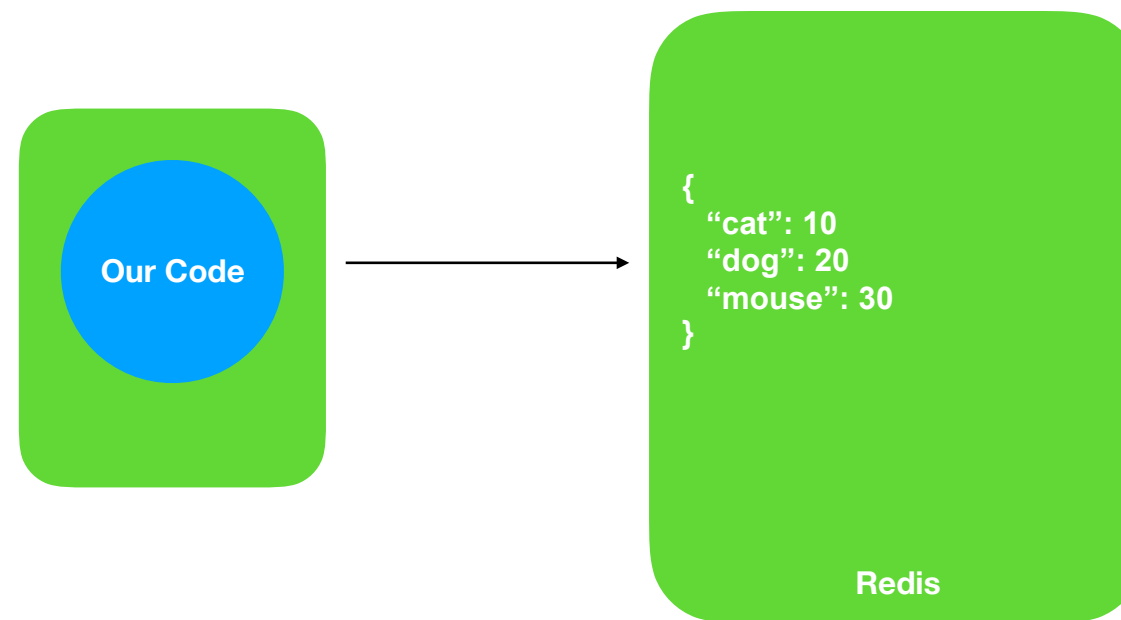
# Redis (cache)



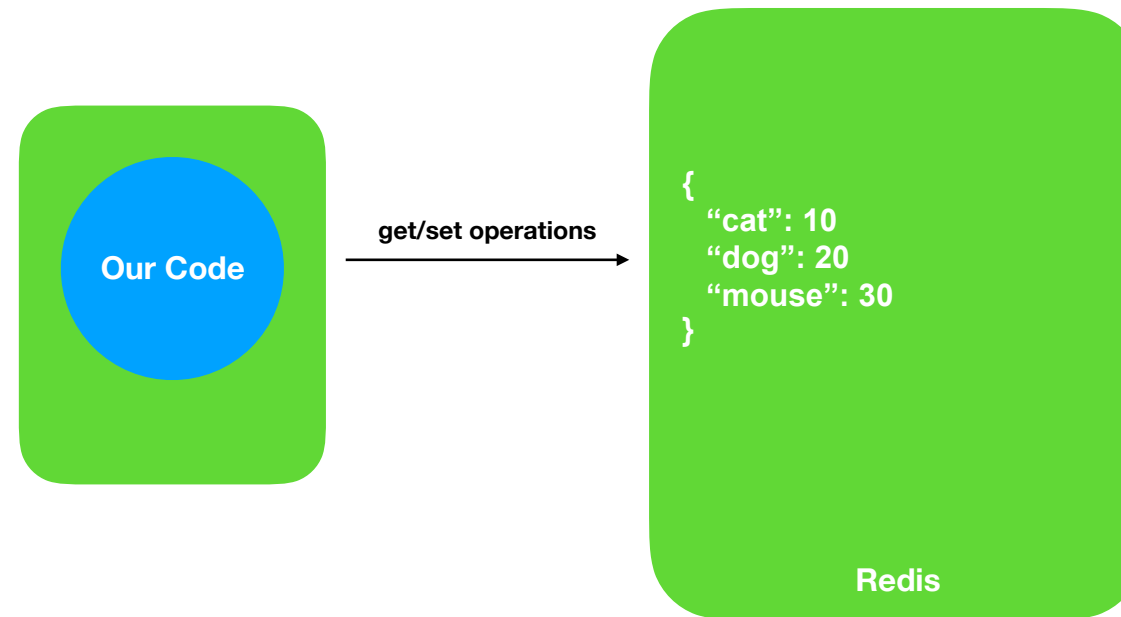
# Redis (cache)



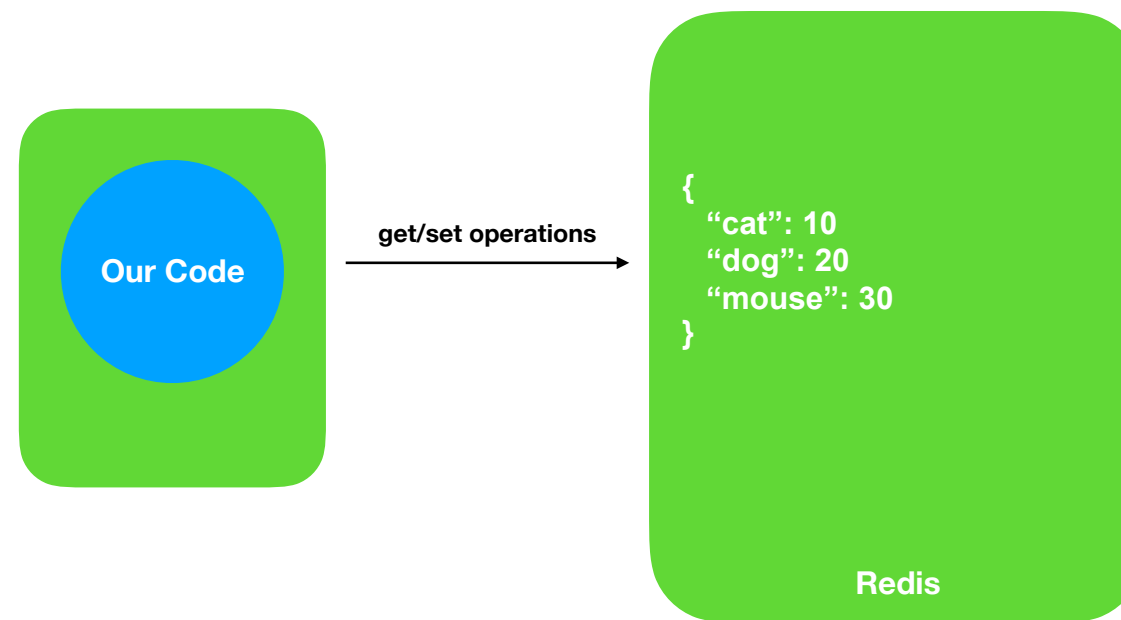
# Redis (cache)



# Redis (cache)




# Redis (cache)



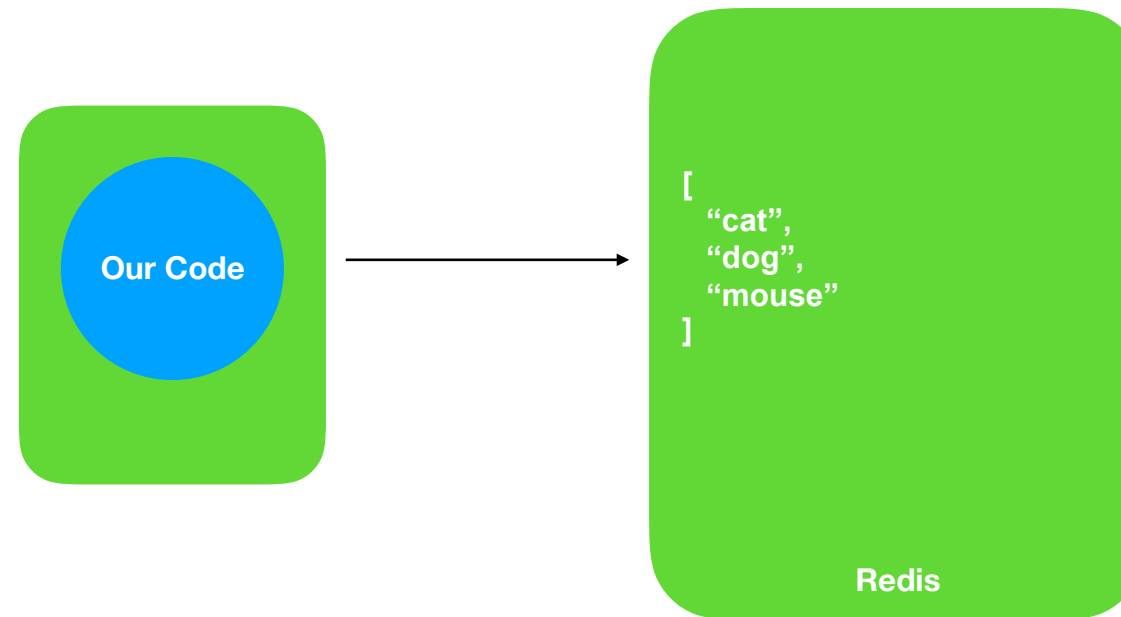
**It's just like using a dictionary on our own machine!**  
**(but its shared and replicated)**



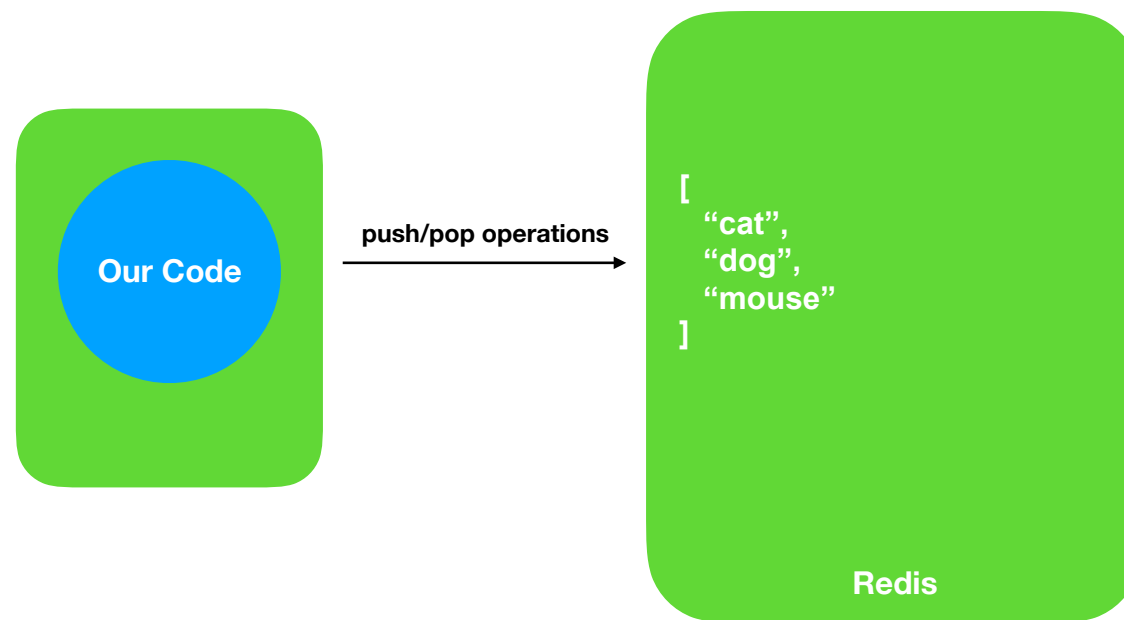
# Redis (lists)



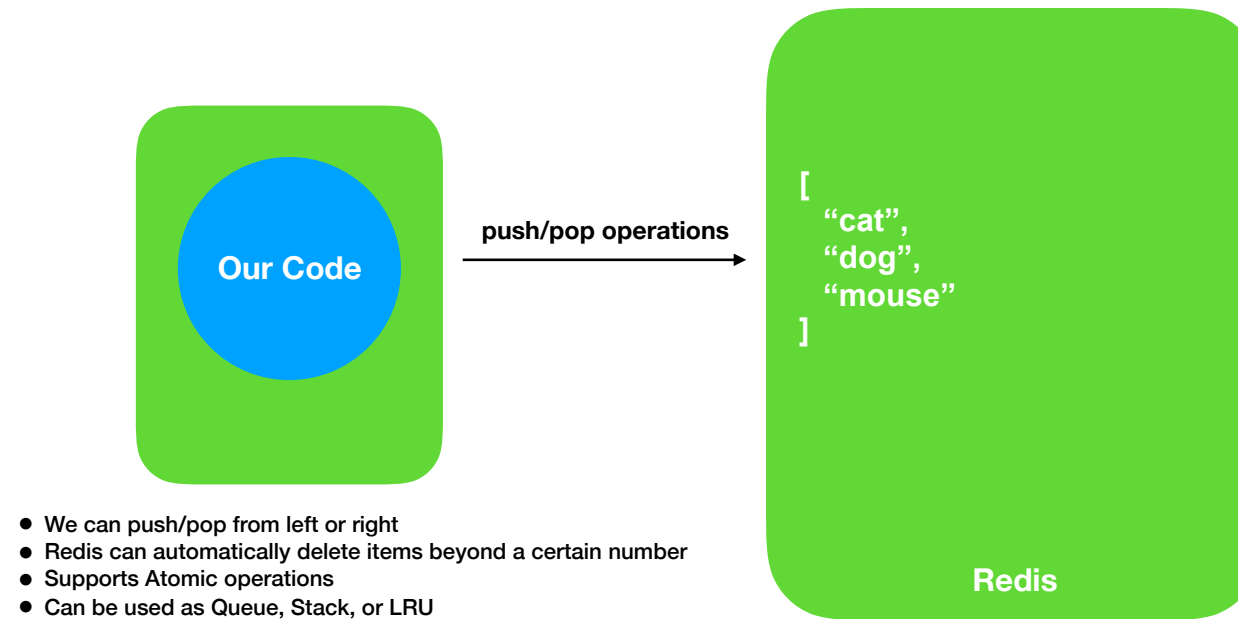
# Redis (lists)



# Redis (lists)



# Redis (lists)




**We can use a queue to distribute work between machines**

# Event-Driven Architecture

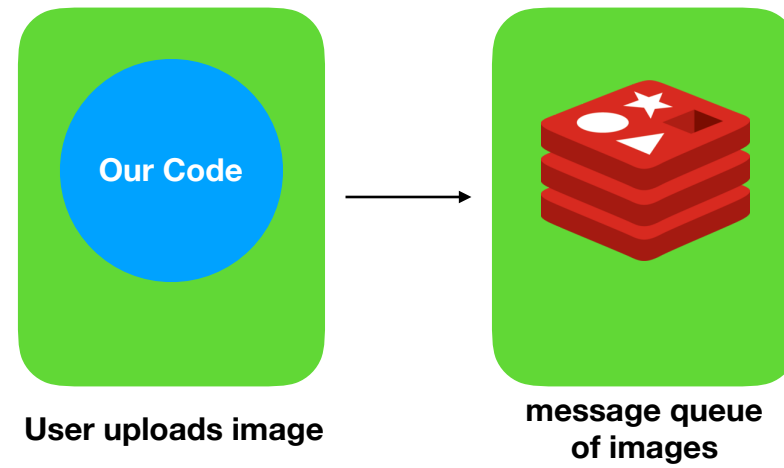
- Ask about the pros/cons of doing it this way, vs. Web server generating images

# Event-Driven Architecture

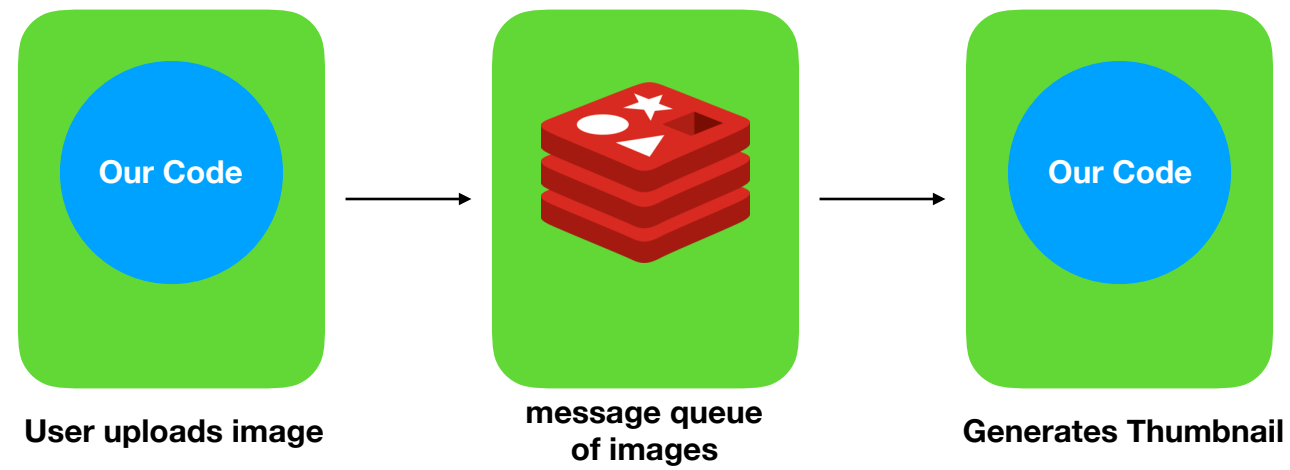


User uploads image

# Event-Driven Architecture

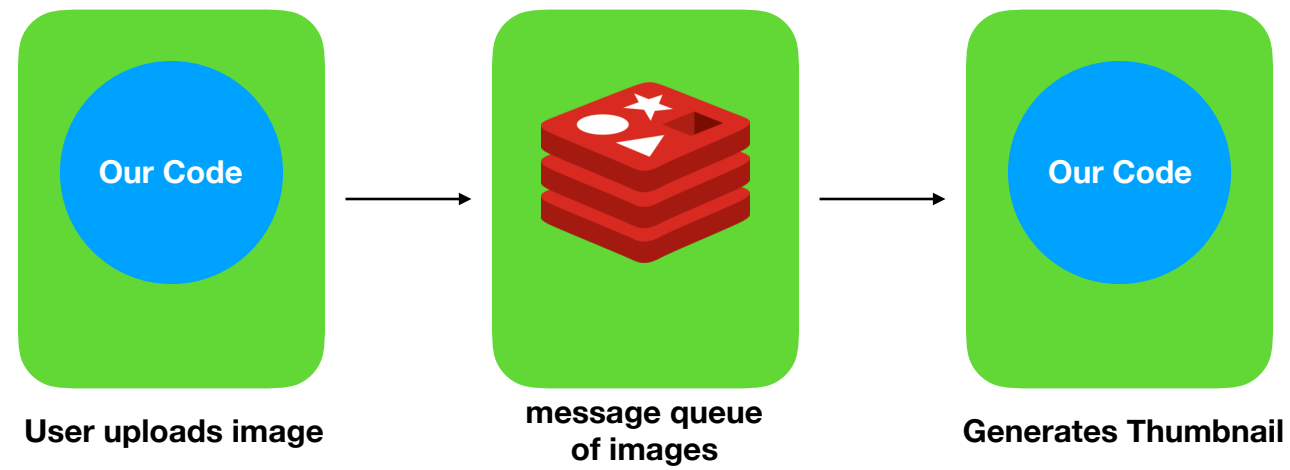


# Event-Driven Architecture





# Event-Driven Architecture




- Our endpoint (uploading images) does less work
- Return to client faster
- More transactions/second for users
- Could scale the 'generate thumbnail' code

# Umbrella Example

Weather service

# Weather Service



- Do a demo! - PyCharm Presentation Mode?

This page lists code related to Protocol Buffers which is developed and maintained by third parties. You may find this code useful, but note that **these projects are not affiliated with or endorsed by Google (unless explicitly marked)**; try them at your own risk. Also note that many projects here are in the early stages of development and not production-ready.

## Programming Languages

- Action Script: <http://code.google.com/p/protobuf-actionscript/>
- Action Script: <https://code.google.com/p/protoc-gen-as3/>
- Action Script: <https://github.com/matrix3d/JProtoc>
- Action Script: <https://github.com/zhongfq/protobuf-as3/>
- C: <https://github.com/protobuf-c/protobuf-c>
- C: <http://koti.kapsi.fi/jpa/nanopb/>
- C: <https://github.com/cloudwu/pbc/>
- C: <https://github.com/haberman/upb/wiki>
- C: <https://github.com/squidfunk/protobluuff>

- [Eclipse editor for protobuf \(from Google\)](#)
- [C++ Builder compatible protobuf](#)
- [Maven Protobuf Compiler Plugin](#)
  - By xolstice.org ([Documentation](#)) ([Source](#)) maven-central v0.6.1
  - <http://igor-petruk.github.com/protobuf-maven-plugin/>
  - <http://code.google.com/p/maven-protoc-plugin/>
  - <https://github.com/os72/protoc-jar-maven-plugin>
- [Documentation generator plugin \(Markdown/HTML/DocBook/...\)](#)
- [DocBook generator for .proto files](#)
- [Protobuf for nginx module](#)
- [RSpec matchers and Cucumber step defs for testing Protocol Buffers](#)
- [Sbt plugin for Protocol Buffers](#)
- [Gradle Protobuf Plugin](#)
- [Multi-platform executable JAR and Java API for protoc](#)
- [Python scripts to convert between Protocol Buffers and JSON](#)
- [Visual Studio Language Service support for Protocol Buffers](#)
- [Visual Studio Code Support for Protocol Buffers](#)
- [C++ library for serialization/de-serialization between Protocol Buffers and JSON.](#)
- [ProtoBuf with Java EE7 Expression Language 3.0; pure Java ProtoBuf Parser and Builder.](#)
- [Notepad++ Syntax Highlighting for .proto files](#)
- [Linter for .proto files](#)
- [Protocol Buffers Dynamic Schema - create protobuf schemas programmatically \(Java\)](#)
- [Make protoc plugins in NodeJS](#)
- [ProfaneDB - A Protocol Buffers database](#)
- [Protocol Buffer property-based testing utility and example message generator \(Python / Hypothesis\)](#)
- [Protolock - CLI utility to prevent backward-incompatible changes to .proto files](#)
- [Optional GRPC - GRPC for testable microservices \(Python\)](#)

- [Eclipse editor for protobuf \(from Google\)](#)
- [C++ Builder compatible protobuf](#)
- [Maven Protobuf Compiler Plugin](#)
  - By xolstice.org ([Documentation](#)) ([Source](#)) maven-central v0.6.1
  - <http://igor-petruk.github.com/protobuf-maven-plugin/>
  - <http://code.google.com/p/maven-protoc-plugin/>
  - <https://github.com/os72/protoc-jar-maven-plugin>
- [Documentation generator plugin \(Markdown/HTML/DocBook/...\)](#)
- [DocBook generator for .proto files](#)
- [Protobuf for nginx module](#)
- [RSpec matchers and Cucumber step defs for testing Protocol Buffers](#)
- [Sbt plugin for Protocol Buffers](#)
- [Gradle Protobuf Plugin](#)
- [Multi-platform executable JAR and Java API for protoc](#)
- [Python scripts to convert between Protocol Buffers and JSON](#)
- [Visual Studio Language Service support for Protocol Buffers](#)
- [Visual Studio Code Support for Protocol Buffers](#)
- [C++ library for serialization/de-serialization between Protocol Buffers and JSON.](#)
- [ProtoBuf with Java EE7 Expression Language 3.0; pure Java ProtoBuf Parser and Builder.](#)
- [Notepad++ Syntax Highlighting for .proto files](#)
- [Linter for .proto files](#)
- [Protocol Buffers Dynamic Schema - create protobuf schemas programmatically \(Java\)](#)
- [Make protoc plugins in NodeJS](#)
- [ProfaneDB - A Protocol Buffers database](#)
- [Protocol Buffer property-based testing utility and example message generator \(Python / Hypothesis\)](#)
- [Protolock - CLI utility to prevent backward-incompatible changes to .proto files](#)
- [Optional GRPC - GRPC for testable microservices \(Python\)](#)



# Optional-GRPC

As a wrapper - if you want less code

# OptionalGRPC (server)

```
@Service(rpc_servicer = add_FooServicer_to_server, stub = FooStub, port = 1000)
class Foo(foo_pb2_grpc.FooServicer):

    def __init__(self, configs: Dict[str, Union[int, str]],
                 server: bool = False,
                 use_rpc: bool = False):
        self.configs = configs
        self.server = server
        self.use_rpc = use_rpc

    /.../

if __name__ == "__main__":
    Foo(configs = {/.../}, server = True, use_rpc = not IS_RUNNING_LOCAL)
```

- Save yourself the boilerplate code for the server!



# OptionalGRPC (client)

```
client = Foo(configs = configs,  
             server = False,  
             use_rpc = not IS_RUNNING_LOCAL)  
resp: foo_pb2.MyMessage = client.do_a_thing(request = MyMessage(text = "hello"))
```

- Save yourself the boilerplate code for the server!

# OptionalGRPC (setup.py)

```
class BuildCommand(build):
    def run(self):
        import optionalgrpc.setup
        optionalgrpc.setup.compile_proto(project_root = "my_foo_project")
        build.run(self)

class InstallCommand(install):
    def run(self):
        import optionalgrpc.setup
        optionalgrpc.setup.compile_proto(project_root = "my_foo_project")
        if not self._called_from_setup(inspect.currentframe()):
            # Run in backward-compatibility mode to support bdist_* commands.
            install.run(self)
        else:
            install.do_egg_install(self) # OR: install.do_egg_install(self)
```

- Compile protobufs as part of pip install!

# OptionalGRPC

- Minimize gRPC Boilerplate
- Test your code as a single process or RPC!
- Compile your code as part of pip install (see examples)
- <https://github.com/mattpaletta/optional-grpc>
- Also have a similar wrapper for Thrift (link on Github)

- Save yourself the boilerplate code for the server!

[github.com/mattpaletta/distributed-docker-example](https://github.com/mattpaletta/distributed-docker-example)

- <https://docs.docker.com/engine/reference/builder/>
- <https://docs.docker.com/compose/compose-file/>
- [https://docs.docker.com/develop/develop-images/dockerfile\\_best-practices/](https://docs.docker.com/develop/develop-images/dockerfile_best-practices/)
- <https://training.play-with-docker.com/alacart/>

# Your Turn!

Questions, Comments, Concerns, Queries, Qwibbles?